



Chapter 4: Threads & Concurrency





Outline

- Overview
- Multicore Programming
- Multithreading Models
- Thread Libraries
- Implicit Threading
- Threading Issues
- Operating System Examples





Objectives

- Identify the basic components of a thread, and contrast threads and processes
- Describe the benefits and challenges of designing multithreaded applications
- Illustrate different approaches to implicit threading including thread pools, fork-join, and Grand Central Dispatch
- Describe how the Windows and Linux operating systems represent threads
- Designing multithreaded applications using the Pthreads, Java, and Windows threading APIs





Motivation

- Most modern applications are multithreaded
- Threads run within application
- Multiple tasks with the application can be implemented by separate threads
 - Update display
 - Fetch data
 - Spell checking
 - Answer a network request
- Process creation is heavy-weight while thread creation is light-weight
- Can simplify code, increase efficiency
- Kernels are generally multithreaded





Aplicatii concurente

- proces = un singur punct de executie in aplicatie (o singura instanta in rulare a aplicatiei)
- unele aplicatii pot profita de existenta mai multor puncte de executie simultana in cadrul aplicatiei (mai ales pe multiprocesoare)
- ex. 1: procesor cu 4 core-uri + aplicatie de filtrare de imagini care imparte imaginea in 4 cadrane, fiecare core filtrand un cadran
 - avantaj: descompunerea activitatilor mari in activitati mai simple care ruleaza simultan => reducerea timpului de rulare
- ex. 2: un singur procesor + program de gestiune a ferestrelor in GUI
 - verifica si proceseaza input-ul de la tastatura, mouse si retea pt. a produce output pe ecran
 - activitati concurente: verificare/procesare input tastatura, mouse si retea, afisare bitmap-uri pe ecran
 - ▶ activitati codate in module separate, cu date private, comunica intre ele prin date partajate






Aplicatii concurente (cont.)

- ex. 3: server de retea
 - activitatea principala: asteapta cereri de la client, le proceseaza si trimite inapoi raspunsurile
 - cu un singur punct de executie procesarea unei cereri particulare ia mult timp (eg, in asteptarea datelor de pe disc)
 - => intarzieri mari pt tratarea altor cereri (situatie similara cu cea care a condus la nevoia de multitasking)
 - solutie: fiecare cerere client e procesata in alt punct de executie al aplicatiei server





Multithreaded Server Architecture





Procese si date partajate

- crearea de puncte multiple de executie in program cu procese care partajeaza date (eg, prin memorie partajata, cu acces protejat cu semafoare/locks)
 - procesele au date private (datorita protectiei MMU a spatiului de adresa)
 - comunica intre ele prin IPC (memorie partajata)
 - daca exista capacitate de multiprocessing, procesele pot rula simultan pe mai multe procesoare
- puncte nevralgice
 - crearea proceselor
 - context-switch-ul
 - IPC-ul





Crearea proceselor

- contextul unui proces e stufos (stocat in PCB in kernel)
 - starea completa a CPU
 - registre de gestiunea memoriei
 - tabele de pagini de memorie
 - descriptori de fisiere
 - actiuni asociate semnalelor, etc

=> cost semnificativ de creare a unui proces
- daca acest cost e prea mare => procesoare folosite ineficient
 - ex aplicatie de filtrare imagini
 - ▶ pt timp de creare proces comparabil cu timpul de filtrare al unui cadran, filtrarea intregii imagini cu 4 procese poate dura mai mult decat filtrarea cu un singur proces (*slowdown*)





Context-switch

- salvarea contextul unui proces + incarcarea contextului unui nou proces => un cost semnificativ avand in vedere bogatia de informatii din PCB
- ex: secventa de evenimente pt. producator-consumator cu zero buffering (*rendezvous*)
 - producatorul produce un element si se blocheaza
 - sistemul face context-switch si aduce consumatorul pe procesor
 - consumatorul consuma elementul si se blocheaza
 - sistemul face context-switch si aduce producatorul inapoi pe procesor pt a produce un nou element
- timp de context-switch mare => viteza de transfer a datelor intre producator si consumator serios afectata
 - fiecare transfer implica doua context-switch-uri
 - oricum, imposibil sa se atinga viteza teoretica maxima de transfer







Fire de executie (threads of execution)

- modalitate de a reduce costuri pentru
 - crearea punctelor de executie in aplicatie
 - schimbarii contextelor de executie intre ele
 - partajarea datelor (IPC)
 - idee: puncte de executie multiple din aplicatie partajeaza o parte din contextul programului (registrele de gestiune a memoriei, tabelele de pagini, descriptorii de fisiere deschise, samd)
 - DAR, fiecare punct de executie are:
 - o copie individuala a unui subset al contextului de executie a aplicatiei (de ex. contextul CPU)
 - structuri de date necesare punctului de executie (de ex. stiva)
- => o noua abstractie de nivel inalt, *firul de executie (thread)* = un punct de executie cu context redus in cadrul programului
- referit uneori din cauza acestei viziuni ca fiind un *proces usor* (lightweight process, LWP)






Single and Multithreaded Processes



single-threaded process



multithreaded process





Benefits

- **Responsiveness** – may allow continued execution if part of process is blocked, especially important for user interfaces
- **Resource Sharing** – threads share resources of process, easier than shared memory or message passing
- **Economy** – cheaper than process creation, thread switching lower overhead than context switching
- **Scalability** – process can take advantage of multicore architectures





Caracteristici threaduri

- ruleaza secvential, au program counter si stiva proprii
- multiplexeaza accesul la CPU ca si procesele (pe multiprocesoare ruleaza in paralel)
- pot crea alte threaduri
- pot executa apeluri de sistem
 - thread blocat in apel sistem => alt thread primeste procesorul
- analogie posibila
 - threadul este fata de proces ceea ce procesul e in raport cu procesorul
 - procesul actioneaza ca un procesor virtual pe care ruleaza thread-ul






Diferente fata de procese

- threadurile aceluiasi proces partajeaza spatiul de adresa al procesului (de ex, partajeaza variabilele globale)
 - => un thread poate distruge usor alt thread (nu exista protectia MMU ca in cazul proceselor diferite)
- lipsa protectiei intre threaduri
 - inevitabila prin design
 - nu e necesara (threadurile sunt parte a aceluiasi program al unui anumit utilizator)
 - nici nu e de dorit (impunerea unor domenii de protectie conduce la probleme similare proceselor)
- alte resurse partajate: acelasi set de fisiere deschise, timere, semnale, etc






Alte caracteristici ale threadurilor

- stari (la fel ca la procese): *in rulare, gata de rulare, blocat, terminat*
 - modele de utilizare
 - cooperativ, lucru in echipa (ex filtrare de imagini)
 - master-slave/worker (ex server)
 - pipeline (ex producator-consumator)
 - avantaj principal: datele partajate sunt datele globale din proces
 - nu e nevoie de mecanisme IPC de tipul memoriei partajate ca in cazul proceselor
 - ▶ ex: buffer global pt. producator-consumator
 - argument puternic pt sistemele multiprocesor unde threadurile pot rula pe CPU-uri diferite
- => partajarea implicita a datelor prin spatiul comun de adresa





Design-ul pachetelor de threaduri

- pachet de threaduri = colectie de primitive (apeluri de biblioteca) pt lucrul cu threaduri
- (1) gestiunea threadurilor
 - *create thread*:
 - ▶ argumente:
 - functia care reprezinta punctul de executie initial
 - stiva privata
 - prioritate de planificare pe procesor
 - ▶ intoarce un TID (Thread ID)
 - *terminare thread*: explicit prin apel *thread exit* sau semnal de tip *kill* de la alt thread/proces
 - primitive de sincronizarea accesului la date partajate
 - ▶ mutex-uri + variabile de conditie
 - ▶ ex: *acquire/release* resurse folosind mutex+variabile conditie






Design-ul pachetelor de threaduri (cont.)

- (2) planificare (aceeasi algoritmi ca si la procese, vom discuta la planificarea proceselor/threadurilor)
 - (3) probleme de reentranta
 - scenariu: doua threaduri T1 si T2 executa concurent apeluri sistem, T1 reuseste, T2 esueaza
 - daca T1 nu evalueaza *errno* inainte ca T2 sa execute apelul sistem care esueaza, T1 va crede eronat ca apelul sau sistem a esuat
 - problema principiala: *errno* e variabila globala
 - solutii posibile:
 - (a) protejarea *errno* cu mutex-uri
 - (b) copie privata a lui *errno* stocata in variabila globala thread-ului apelant, dar privata (invizibila) celorlalte thread-uri
- => TLS (Thread-Local Storage)






Implementarea threadurilor kernel

- pachetele de threaduri se pot implementa in kernel sau in spatiul utilizator
- threadurile kernel separa campurile din PCB care ajuta la crearea unui punct de executie si le stocheaza intr-un TCB
 - => un proces cu un singur punct de executie e reprezentat **in kernel** de un PCB si un TCB
- operatia de creare a unui thread (*thread_create*) este un apel sistem
 - aloca un TCB
 - aloca stive kernel si user
 - le leaga la PCB-ul procesului in care s-a facut apelul sistem





Exemplu cu doua threaduri kernel

- un thread suspendat in kernel
- altul ruland in spatiul utilizator
- pt ca sunt implementate in kernel si seamana cu procesele, threadurile kernel se mai cheama si *procesuse usoare* (*lightweight processes, LWP*)




Figure 3: Example of kernel thread implementation.





Costuri threaduri kernel

- cost creare thread kernel $\ll$ cost creare proces
 - se alocă și initializează doar TCB-ul și stivele
 - restul contextului există deja creat în PCB
- context switch-ul
 - între două threaduri **ale aceluiași proces**: mai rapid decât context switch-ul a două procese (nu trebuie schimbat restul contextului din PCB)
 - între două threaduri din procese **diferite**: același cost ca și context switch-ul de procese (nu se schimbă doar TCB-urile, ci și PCB-urile)





Planificarea kernel threadurilor kernel

- threadurile ruleaza asincron unele fata de celelalte si pot pierde procesorul la fel ca si procesele
 - => accesul la date partajate trebuie sincronizat
- planificatorul kernel alege urmatorul thread care trebuie sa ruleze
 - => politicile de planificare ale aplicatiei (de ex bazata pe prioritati) trebuie comunicate intr-un fel sau altul kernelului
- pt. multiprocesoare, planificatorul poate aloca mai multe CPU-uri unui singur proces pt. ca threadurile sa ruleze in paralel (*gang scheduling*)
- thread blocat in kernel (de ex prin apel de sistem I/O) + cuanta de timp proces neconsumata integral
 - => planificatorul cauta in lista de TCB-uri un thread gata de rulare **din acelasi proces** si ii acorda procesorul
 - consecinta: lightweight context switch, TCB vs PCB





Protectia threadurilor kernel

- observatia generala despre threaduri e valabila si pt threaduri kernel
- un thread poate corupe stiva altui thread, de pilda
 - => se distrug datele private ale altui thread (variabilele automate/locale)
- solutie: implementarea stivelor in spatii de adresa diferite, dar asta mareste costul context switch-ului
 - necesita salvarea si reincarcarea contextelor de executie referitoare la gestiunea memoriei, pe langa contextul uzual din TCB





Dezavantajele threadurilor kernel

- desi mai putin costisitoare ca procesele, anumite aspecte le fac nepotrivite pt utilizatori
 - (1) *thread_create* este apel sistem => costisitor pt procese care creeaza multe threaduri
 - (2) context switching-ul de threaduri necesita intrarea si iesirea in/din kernel mode => overhead aditional context switching-ului obisnuit
 - (3) implementarea in kernel e **inflexibila**
 - impune un model de threaduri care nu e potrivit pt orice aplicatie
 - codul planificatorului (scheduler) nu e accesibil (fiind in kernel)
- => greu de adaptat pt cerintele specifice ale unei anumite aplicatii (politica de planificare nu se poate schimba usor)





User-level threads

- daca planificatorul si TCB-urile sunt implementate in spatiul utilizator costurile scad pentru ca nu mai e nevoie de transgresarea granitei kernel/user
- comparatie calitativa, intr-un sistem in care un apel de procedura costa 7 usec, iar un apel sistem (trap) 19 usec




Figure 4: Example of user-level thread implementation.

Operatie	Thread user	Thread kernel	Proces Unix
fork	34 usec	948 usec	11300 usec
IPC synch	37 usec	441 usec	1840 usec





Caracteristici threaduri utilizator

- nu apeleaza serviciile kernel pt creare si context switch => aceste operatii sunt f. rapide
 - nu se trece granita user-kernel si nu e nevoie de verificarea parametrilor apelului sistem de catre kernel
- aplicatiile pot furniza propriul planificator (*thread scheduler*) customizat cf unor cerinte specifice
- nu necesita suport explicit din partea kernelului
 - procesul e privit ca un procesor virtual pt. threaduri
- fiind mult mai rapide decat threadurile kernel, de regula threadurile user se implementeaza deasupra threadurilor kernel
 - => planificatorul de threaduri user trateaza threadurile kernel ca pe procesoare virtuale
 - ▶ multiplexeaza mai multe threaduri user pe unul sau mai multe threaduri kernel






User Threads and Kernel Threads

- **User threads** - management done by user-level threads library
- Three primary thread libraries:
 - POSIX **Pthreads**
 - Windows threads
 - Java threads
- **Kernel threads** - Supported by the Kernel
- Examples – virtually all general-purpose operating systems, including:
 - Windows
 - Linux
 - Mac OS X
 - iOS
 - Android





User and Kernel Threads





Multithreading Models


- Many-to-One
- One-to-One
- Many-to-Many





Many-to-One


- Many user-level threads mapped to single kernel thread
- One thread blocking causes all to block
- Multiple threads may not run in parallel on multicore system because only one may be in kernel at a time
- Few systems currently use this model
- Examples:
 - **Solaris Green Threads**
 - **GNU Portable Threads**





One-to-One


- Each user-level thread maps to kernel thread
- Creating a user-level thread creates a kernel thread
- More concurrency than many-to-one
- Number of threads per process sometimes restricted due to overhead
- Examples
 - Windows
 - Linux





Many-to-Many Model


- Allows many user level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads
- Windows with the *ThreadFiber* package
- Otherwise not very common





Two-level Model

- Similar to M:M, except that it allows a user thread to be **bound** to kernel thread





Probleme comune threadurilor k. si u.

- reentranta: multe apeluri de biblioteca sunt ne-reentrante
 - apelurile reentrante desemnate in manual ca “thread safe”
 - ex: trimiterea unui mesaj in retea in 2 pasi: (a) asamblare mesaj in buffer + (b) transmisie (*syscall*)
 - pierderea procesorului intre (a) si (b) poate conduce la suprascrierea bufferului
 - alte ex: *errno*, *malloc*, apeluri *stdio*
- tratarea semnalelor
 - ex: un thread trateaza un semnal in vreme ce alt thread vrea ca semnalul respectiv sa termine aplicatia (procesul, mai exact)
 - se poate intampla daca se folosesc apeluri de biblioteca si runtime user impreuna
 - cauza: semnalele sunt definite per proces si nu per thread





Dezavantaje threaduri utilizator

- determinate de faptul ca existenta lor e necunoscuta kernelului
- (1) thread user executa apel sistem blocant in kernel
 - planificatorul kernel (agnostic cu privire la threadurile user) considera intreg procesul blocat
 - se alocă CPU altui proces (sau kernel thread), *chiar daca procesul curent mai are si alte threaduri user gata de rulare si nu si-a consumat toata cuanta de timp CPU !*
- (2) thread user comite *page fault* (acces la pagina de memorie inexistentă/nealocată)
 - planificatorul alege alt proces/kernel thread pt executie pe durata transferului paginii de pe disc in memorie
 - => aplicatia ruleaza cu mai putine CPU decat e necesar
- (3) kernelul scoate de pe CPU un thread user care detine un (spin)lock ne-eliberat
 - => scade dramatic performanta aplicatiilor care folosesc threaduri user in paralel
 - efect exacerbat daca sunt alese sa ruleze alte threaduri care vor sa obtina acelasi spinlock
- *problema de fond*: lipsa de coordonare scheduler kernel <—> pachet de threaduri user





Scheduler activations

- metoda de coordonare kernel $\leftrightarrow$ pachet de threaduri utilizator
- model: aplicatia ruleaza pe un multiprocesor virtual
 - apeluri sistem pt a suplimenta/diminua nr de procesoare virtuale alocate de kernel (“adauga CPU”, “acest CPU este idle”)
 - kernelul decide daca onoreaza cererea sau nu
- scheduler activation e aproximarea unui kernel thread
 - are stive kernel si user
 - ofera context de executie pt un thread user
 - in plus, ofera conceptul de *upcall* pt evenimente de mai multe tipuri
 - ▶ fiecare upcall creeaza o noua activare (optimizare: folosirea vechilor activari)





Tipuri de evenimente upcall

- *adauga procesor* – consecinta: executia unui thread user
- *procesor preemptat* – adauga threadul user care se executa in activarea care a pierdut CPU in coada de threaduri gata de rulare
- *activare blocata* – activarea s-a blocat (eg, apel sistem blocant) si nu mai utilizeaza procesorul
- *activare deblocata* – pune in lista gata de rulare threadul care se executa in contextul activarii blocate





Functionarea scheduler activations

- la pornirea procesului
 - kernelul alocă o activare + notifica aplicația (*upcall* “adauga CPU”) după ce i-a asigurat un CPU
 - sistemul de gestiune al threadurilor user primește notificarea
 - ▶ folosește noua activare drept context de execuție pt initializarea sa și a threadului *main*
 - ▶ ulterior se pot crea noi threaduri și cere noi procesoare
- la cerere de suplimentare a concurenței (adaugare CPU sau pt. thread blocat)
 - kernelul salvează starea threadului user în activarea curentă
 - alocă o nouă activare și cheamă aplicația în contextul noii activări





Functionarea scheduler activations (cont)


- la blocarea unui thread user in kernel
 - nu se continua in kernel la deblocare
 - se creeaza o noua activare
 - se notifica aplicatia
 - schedulerul user copiaza starea threadului blocat din vechea activare si notifica kernelul ca vechea activare poate fi refolosita
- cand un thread user care detine un spinlock pierde procesorul
 - kernelul genereaza un *upcall*
 - sistemul de gestiune a threadurilor verifica daca threadul e in sectiune critica
 - in caz afirmativ, threadul este continuat temporar printr-un context switch in user space
 - la iesirea din sectiunea critica, se da controlul inapoi *upcall-ului* tot prin context switch in user space





Scheduler Activations

- Both M:M and Two-level models require communication to maintain the appropriate number of kernel threads allocated to the application
- Typically use an intermediate data structure between user and kernel threads – **lightweight process (LWP)**
 - Appears to be a virtual processor on which process can schedule user thread to run
 - Each LWP attached to kernel thread
 - How many LWPs to create?
- Scheduler activations provide **upcalls** - a communication mechanism from the kernel to the **upcall handler** in the thread library
- This communication allows an application to maintain the correct number kernel threads





Thread Libraries

- **Thread library** provides programmer with API for creating and managing threads
- Two primary ways of implementing
 - Library entirely in user space
 - Kernel-level library supported by the OS





Exemple de pachete de threaduri user

- Solaris si POSIX
- threaduri Solaris
 - exista si la nivel de kernel (s.n. LWPs, sunt multiplexate pe CPU existente) si la nivel de utilizator
 - pachetul de threaduri (*libthread*) planifica threadurile user pe o colectie de LWPs
 - uzual, threadurile user sunt create nelegate (*unbound*), adica se pot muta intre LWP-uri si asa se si folosesc in general
 - daca un thread user e asignat unui LWP se spune ca e legat (*bound*)
 - controlul kernelului asupra threadurilor user se exercita prin bound threads cu ajutorul setarii prioritatii LWP-ului asociat (de ex, thread user pt stream audio cu prioritate mare in aplicatie multimedia)





Thread-uri Solaris (cont.)



....

- biblioteca de threaduri asigura ca exista suficiente LWP-uri active pt ca procesul sa poata continua
- daca un proces determina toate LWP-urile sale sa se blocheze, biblioteca va crea LWP-uri aditionale pt. ca threadurile neblocate sa poata rula
- daca LWP-urile sunt idle prea mult timp, biblioteca le da inapoi kernelului pt a fi folosite de alte aplicatii





Thread-uri Solaris (cont.)

- exista apel de setare explicita a nr de LWP-uri la un nou nivel
thr_setconcurrency(int new_level)
- mecanisme IPC: mutex si variabile conditie
- implementeaza TLS (Thread-Local Storage), capacitate privata de stocare a variabilelor in afara stivei
 - variabile locale nepartajate, “statice” (globale pt threadul respectiv)
 - declarate cu #pragma, de ex #pragma unshared errno;





Solaris threads

```
int thr_create(void *stack_base, size_t stack_size,  
              void *(*function)(void *), void *arg, long flags)  
void thr_exit(void *status)  
int thr_join(thread_t wait_for, thread_t *joiner_id, void *status)  
int mutex_init(mutex_t *mutex, int type, void *arg)  
int mutex_lock(mutex_t *mutex)  
int mutex_unlock(mutex_t *mutex);  
int cond_init(cond_t *cond, int type, int arg)  
int cond_wait(cond_t *cond, mutex_t *mutex)  
int cond_signal(cond_t *cond)  
int cond_broadcast(cond_t *cond)
```

Figure 5: Some commonly used UI thread operations.





Tratarea semnalelor

- fiecare thread are propria masca de semnale
- toate threadurile unui proces partajeaza handlerele de semnal
- cand SIG_IGN/SIG_DFL sunt setate, se aplica tuturor threadurilor procesului
- semnalele de tip *trap* (sincrone cu executia threadului, ex SIGFPE, SIGSEGV) executate exclusiv de threadul care le-a generat
 - => m. m. threaduri pot genera si trata acelasi tip de semnal simultan
- semnalele de tip *intrerupere* (generate asincron cu executia threadului, ex SIGIO, SIGINT)
 - pot fi tratate de orice thread care nu mascheaza/blocheaza semnalul respectiv
 - dc. sunt mai multe astfel de threaduri, se alege unul dintre ele pt tratarea semnalului
- daca toate threadurile au mascat un anume semnal, se asteapta dupa primul thread care deblocheaza tratarea semnalului respectiv





Tratarea semnalelor (cont.)

- *thread_kill*: trimite un semnal catre un thread din acelasi proces
 - semnalul e de tip trap, tratat doar de threadul caruia ii este adresat
 - nu se pot trimite semnale unui *anume* thread dintr-un *alt* proces
- threadurile Solaris implementeaza un nou semnal, SIGWAITING (SIG_IGN implicit)
 - trimis procesului cand toate LWP-urile sunt blocate indefinit in asteptarea unui eveniment extern
 - poate fi folosit pt crearea de noi LWP-uri pt a evita blocarea intregului proces
 - idee similara cu scheduler activations, dar coarse-grain (nu merge pt short-term blocking, doar pt evenimente de tip page faults, filesystem I/O)





Thread-uri POSIX

- *pthread*
- standard IEEE, Portable Operating System Interface (e o specificatie, nu o implementare)
- asemanatoare threadurilor Solaris





Pthreads

- May be provided either as user-level or kernel-level
- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- ***Specification***, not ***implementation***
- API specifies behavior of the thread library, implementation is up to development of the library
- Common in UNIX operating systems (Linux & Mac OS X)





POSIX threads

```
int pthread_create(pthread_t *thread_pointer,  
                  pthread_attr_t *attributes,  
                  void *(*function)(void *), void *arg)  
void pthread_exit(void *status)  
int pthread_join(pthread_t thread, void **status)  
void pthread_mutex_init(pthread_mutex_t *mutex,  
                        pthread_mutexattr_t *attr)  
int pthread_mutex_lock(pthread_mutex_t *mutex)  
int pthread_mutex_unlock(pthread_mutex_t *mutex)  
int pthread_cond_init(pthread_cond_t *cond, pthread_condattr_t *attr)  
int pthread_cond_wait(pthread_cond_t *cond, pthread_mutex_t *mutex)  
int pthread_cond_signal(pthread_cond_t *cond)  
int pthread_cond_broadcast(pthread_cond_t *cond)
```

Figure 6: Some commonly used Pthreads operations.





Pthreads Example

```
#include <pthread.h>
#include <stdio.h>

#include <stdlib.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    /* set the default attributes of the thread */
    pthread_attr_init(&attr);
    /* create the thread */
    pthread_create(&tid, &attr, runner, argv[1]);
    /* wait for the thread to exit */
    pthread_join(tid, NULL);

    printf("sum = %d\n", sum);
}
```






Pthreads Example (Cont.)

```
/* The thread will execute in this function */  
void *runner(void *param)  
{  
    int i, upper = atoi(param);  
    sum = 0;  
  
    for (i = 1; i <= upper; i++)  
        sum += i;  
  
    pthread_exit(0);  
}
```





Pthreads Code for Joining 10 Threads

```
#define NUM_THREADS 10

/* an array of threads to be joined upon */
pthread_t workers[NUM_THREADS];

for (int i = 0; i < NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```





Implicit Threading

- Growing in popularity as numbers of threads increase, program correctness more difficult with explicit threads
- Creation and management of threads done by compilers and run-time libraries rather than programmers
- Five methods explored
 - Thread Pools
 - Fork-Join
 - OpenMP
 - Grand Central Dispatch
 - Intel Threading Building Blocks





Thread Pools

- Create a number of threads in a pool where they await work
- Advantages:
 - Usually slightly faster to service a request with an existing thread than create a new thread
 - Allows the number of threads in the application(s) to be bound to the size of the pool
 - Separating task to be performed from mechanics of creating task allows different strategies for running task
 - ▶ i.e, Tasks could be scheduled to run periodically
- Windows API supports thread pools:


```
DWORD WINAPI PoolFunction(AVOID Param) {  
    /*  
    * this function runs as a separate thread.  
    */  
}
```





Fork-Join Parallelism

- Multiple threads (tasks) are **forked**, and then **joined**.





Fork-Join Parallelism

- General algorithm for fork-join strategy:

```
Task(problem)
  if problem is small enough
    solve the problem directly
  else
    subtask1 = fork(new Task(subset of problem))
    subtask2 = fork(new Task(subset of problem))


    result1 = join(subtask1)
    result2 = join(subtask2)


    return combined results
```





Fork-Join Parallelism





Threading Issues

- Semantics of **fork()** and **exec()** system calls
- Signal handling
 - Synchronous and asynchronous
- Thread cancellation of target thread
 - Asynchronous or deferred
- Thread-local storage
- Scheduler Activations





Semantics of `fork()` and `exec()`

- Does `fork()` duplicate only the calling thread or all threads?
 - Some UNIXes have two versions of `fork`
- `exec()` usually works as normal – replace the running process including all threads





Signal Handling

- **Signals** are used in UNIX systems to notify a process that a particular event has occurred.
- A **signal handler** is used to process signals
 1. Signal is generated by particular event
 2. Signal is delivered to a process
 3. Signal is handled by one of two signal handlers:
 1. default
 2. user-defined
- Every signal has **default handler** that kernel runs when handling signal
 - **User-defined signal handler** can override default
 - For single-threaded, signal delivered to process





Signal Handling (Cont.)

- Where should a signal be delivered for multi-threaded?
 - Deliver the signal to the thread to which the signal applies
 - Deliver the signal to every thread in the process
 - Deliver the signal to certain threads in the process
 - Assign a specific thread to receive all signals for the process





Thread Cancellation

- Terminating a thread before it has finished
- Thread to be canceled is **target thread**
- Two general approaches:
 - **Asynchronous cancellation** terminates the target thread immediately
 - **Deferred cancellation** allows the target thread to periodically check if it should be cancelled
- Pthread code to create and cancel a thread:

```
pthread_t tid;

/* create the thread */
pthread_create(&tid, 0, worker, NULL);

. . .

/* cancel the thread */
pthread_cancel(tid);

/* wait for the thread to terminate */
pthread_join(tid, NULL);
```





Thread Cancellation (Cont.)

- Invoking thread cancellation requests cancellation, but actual cancellation depends on thread state

Mode	State	Type
Off	Disabled	–
Deferred	Enabled	Deferred
Asynchronous	Enabled	Asynchronous

- If thread has cancellation disabled, cancellation remains pending until thread enables it
- Default type is deferred
 - Cancellation only occurs when thread reaches **cancellation point**
 - ▶ i.e., `pthread_testcancel()`
 - ▶ Then **cleanup handler** is invoked
- On Linux systems, thread cancellation is handled through signals





Thread-Local Storage

- **Thread-local storage (TLS)** allows each thread to have its own copy of data
- Useful when you do not have control over the thread creation process (i.e., when using a thread pool)
- Different from local variables
 - Local variables visible only during single function invocation
 - TLS visible across function invocations
- Similar to `static` data
 - TLS is unique to each thread





Multicore Programming


- **Multicore** or **multiprocessor** systems puts pressure on programmers, challenges include:
 - **Dividing activities**
 - **Balance**
 - **Data splitting**
 - **Data dependency**
 - **Testing and debugging**
- **Parallelism** implies a system can perform more than one task simultaneously
- **Concurrency** supports more than one task making progress
 - Single processor / core, scheduler providing concurrency






Concurrency vs. Parallelism

- **Concurrent execution on single-core system:**



- **Parallelism on a multi-core system:**






Multicore Programming

- Types of parallelism
 - **Data parallelism** – distributes subsets of the same data across multiple cores, same operation on each
 - **Task parallelism** – distributing threads across cores, each thread performing unique operation





Data and Task Parallelism





Amdahl's Law

- Identifies performance gains from adding additional cores to an application that has both serial and parallel components
- S is serial portion
- N processing cores

$$speedup \leq \frac{1}{S + \frac{(1-S)}{N}}$$

- That is, if application is 75% parallel / 25% serial, moving from 1 to 2 cores results in speedup of 1.6 times
- As N approaches infinity, speedup approaches $1 / S$


Serial portion of an application has disproportionate effect on performance gained by adding additional cores

- But does the law take into account contemporary multicore systems?





Amdahl's Law





Windows Multithreaded C Program

```
#include <windows.h>
#include <stdio.h>
DWORD Sum; /* data is shared by the thread(s) */

/* The thread will execute in this function */
DWORD WINAPI Summation(LPVOID Param)
{
    DWORD Upper = *(DWORD*)Param;
    for (DWORD i = 1; i <= Upper; i++)
        Sum += i;
    return 0;
}
```





Windows Multithreaded C Program (Cont.)

```
int main(int argc, char *argv[])
{
    DWORD ThreadId;
    HANDLE ThreadHandle;
    int Param;

    Param = atoi(argv[1]);
    /* create the thread */
    ThreadHandle = CreateThread(
        NULL, /* default security attributes */
        0, /* default stack size */
        Summation, /* thread function */
        &Param, /* parameter to thread function */
        0, /* default creation flags */
        &ThreadId); /* returns the thread identifier */

    /* now wait for the thread to finish */
    WaitForSingleObject(ThreadHandle, INFINITE);

    /* close the thread handle */
    CloseHandle(ThreadHandle);

    printf("sum = %d\n", Sum);
}
```





Java Threads

- Java threads are managed by the JVM
- Typically implemented using the threads model provided by underlying OS
- Java threads may be created by:
 - Extending Thread class
 - Implementing the Runnable interface

```
public interface Runnable
{
    public abstract void run();
}
```

- Standard practice is to implement Runnable interface





Java Threads

Implementing Runnable interface:

```
class Task implements Runnable
{
    public void run() {
        System.out.println("I am a thread.");
    }
}
```

Creating a thread:

```
Thread worker = new Thread(new Task());
worker.start();
```

Waiting on a thread:

```
try {
    worker.join();
}
catch (InterruptedException ie) { }
```





Java Executor Framework

- Rather than explicitly creating threads, Java also allows thread creation around the Executor interface:

```
public interface Executor
{
    void execute(Runnable command);
}
```

- The Executor is used as follows:

```
Executor service = new Executor();
service.execute(new Task());
```





Java Executor Framework


```
import java.util.concurrent.*;

class Summation implements Callable<Integer>
{
    private int upper;
    public Summation(int upper) {
        this.upper = upper;
    }

    /* The thread will execute in this method */
    public Integer call() {
        int sum = 0;
        for (int i = 1; i <= upper; i++)
            sum += i;

        return new Integer(sum);
    }
}
```





Java Executor Framework (Cont.)

```
public class Driver
{
    public static void main(String[] args) {
        int upper = Integer.parseInt(args[0]);

        ExecutorService pool = Executors.newSingleThreadExecutor();
        Future<Integer> result = pool.submit(new Summation(upper));

        try {
            System.out.println("sum = " + result.get());
        } catch (InterruptedException | ExecutionException ie) { }
    }
}
```





Java Thread Pools

- Three factory methods for creating thread pools in Executors class:
 - `static ExecutorService newSingleThreadExecutor()`
 - `static ExecutorService newFixedThreadPool(int size)`
 - `static ExecutorService newCachedThreadPool()`





Java Thread Pools (Cont.)

```
import java.util.concurrent.*;


public class ThreadPoolExample
{
    public static void main(String[] args) {
        int numTasks = Integer.parseInt(args[0].trim());

        /* Create the thread pool */
        ExecutorService pool = Executors.newCachedThreadPool();

        /* Run each task using a thread in the pool */
        for (int i = 0; i < numTasks; i++)
            pool.execute(new Task());

        /* Shut down the pool once all threads have completed */
        pool.shutdown();
    }
}
```





Fork-Join Parallelism in Java

```
ForkJoinPool pool = new ForkJoinPool();  
// array contains the integers to be summed  
int[] array = new int[SIZE];  
  
SumTask task = new SumTask(0, SIZE - 1, array);  
int sum = pool.invoke(task);
```





Fork-Join Parallelism in Java

```
import java.util.concurrent.*;

public class SumTask extends RecursiveTask<Integer>
{
    static final int THRESHOLD = 1000;

    private int begin;
    private int end;
    private int[] array;

    public SumTask(int begin, int end, int[] array) {
        this.begin = begin;
        this.end = end;
        this.array = array;
    }

    protected Integer compute() {
        if (end - begin < THRESHOLD) {
            int sum = 0;
            for (int i = begin; i <= end; i++)
                sum += array[i];

            return sum;
        }
        else {
            int mid = (begin + end) / 2;

            SumTask leftTask = new SumTask(begin, mid, array);
            SumTask rightTask = new SumTask(mid + 1, end, array);

            leftTask.fork();
            rightTask.fork();


            return rightTask.join() + leftTask.join();
        }
    }
}
```





Fork-Join Parallelism in Java

- The `ForkJoinTask` is an abstract base class
- `RecursiveTask` and `RecursiveAction` classes extend `ForkJoinTask`
- `RecursiveTask` returns a result (via the return value from the `compute()` method)
- `RecursiveAction` does not return a result





OpenMP

- Set of compiler directives and an API for C, C++, FORTRAN
- Provides support for parallel programming in shared-memory environments
- Identifies **parallel regions** – blocks of code that can run in parallel

#pragma omp parallel

Create as many threads as there are cores

```
#include <omp.h>
#include <stdio.h>


int main(int argc, char *argv[])
{
    /* sequential code */

    #pragma omp parallel
    {
        printf("I am a parallel region.");
    }

    /* sequential code */

    return 0;
}
```





- Run the for loop in parallel

```
#pragma omp parallel for  
for (i = 0; i < N; i++) {  
    c[i] = a[i] + b[i];  
}
```





Grand Central Dispatch

- Apple technology for macOS and iOS operating systems
- Extensions to C, C++ and Objective-C languages, API, and run-time library
- Allows identification of parallel sections
- Manages most of the details of threading
- Block is in “`^ { }`” :

```
^{ printf("I am a block"); }
```

- Blocks placed in dispatch queue
 - Assigned to available thread in thread pool when removed from queue





Grand Central Dispatch

- Two types of dispatch queues:
 - **serial** – blocks removed in FIFO order, queue is per process, called **main queue**
 - ▶ Programmers can create additional serial queues within program
 - **concurrent** – removed in FIFO order but several may be removed at a time
 - ▶ Four system wide queues divided by quality of service:
 - `QOS_CLASS_USER_INTERACTIVE`
 - `QOS_CLASS_USER_INITIATED`
 - `QOS_CLASS_USER_UTILITY`
 - `QOS_CLASS_USER_BACKGROUND`





Grand Central Dispatch

- For the Swift language a task is defined as a closure – similar to a block, minus the caret
- Closures are submitted to the queue using the `dispatch_async()` function:

```
let queue = dispatch_get_global_queue  
            (QOS_CLASS_USER_INITIATED, 0)
```

```
dispatch_async(queue, { print("I am a closure.") })
```





Intel Threading Building Blocks (TBB)

- Template library for designing parallel C++ programs
- A serial version of a simple for loop

```
for (int i = 0; i < n; i++) {  
    apply(v[i]);  
}
```

- The same for loop written using TBB with `parallel_for` statement:

```
parallel_for (size_t(0), n, [=](size_t i) {apply(v[i]);});
```





Thread Cancellation in Java

- Deferred cancellation uses the `interrupt()` method, which sets the interrupted status of a thread.

```
Thread worker;
```


```
...
```

```
/* set the interruption status of the thread */  
worker.interrupt()
```

- A thread can then check to see if it has been interrupted:

```
while (!Thread.currentThread().isInterrupted()) {  
    ...  
}
```





Operating System Examples

- Windows Threads
- Linux Threads





Windows Threads

- Windows API – primary API for Windows applications
- Implements the one-to-one mapping, kernel-level
- Each thread contains
 - A thread id
 - Register set representing state of processor
 - Separate user and kernel stacks for when thread runs in user mode or kernel mode
 - Private data storage area used by run-time libraries and dynamic link libraries (DLLs)
- The register set, stacks, and private storage area are known as the **context** of the thread






Windows Threads (Cont.)

- The primary data structures of a thread include:
 - ETHREAD (executive thread block) – includes pointer to process to which thread belongs and to KTHREAD, in kernel space
 - KTHREAD (kernel thread block) – scheduling and synchronization info, kernel-mode stack, pointer to TEB, in kernel space
 - TEB (thread environment block) – thread id, user-mode stack, thread-local storage, in user space





Windows Threads Data Structures





Linux Threads

- Linux refers to them as **tasks** rather than **threads**
- Thread creation is done through `clone()` system call
- `clone()` allows a child task to share the address space of the parent task (process)
 - Flags control behavior

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.

- `struct task_struct` points to process data structures (shared or unique)



End of Chapter 4

